



Leveraging Deep Learning For Emotion Recognition & Detection

Scott Hameed

Springboard Capstone 3





Contents:

- Summary (TL;DR)
- Business Problem
- Scope, Approach, Process, and Libraries
- Neural Network & Algorithms Overview
- Code Overview
- Baseline Performance & Comparison
- Optimization & Results
- Key Challenges & Learnings For Future
- Future Enhancements & Optimizations
- Resources (Source Code & Tech Report)
- Appendix (Helpful Links & Other Resources)



Summary (TL;DR)



- This project leverages Deep Learning to recognize and detect human emotions in images
- Human emotions are fundamental to our experience on our planet, as emotions affect everyone around us [1]
- One of the key inspirations for this project came from interacting with people; I realized that my interaction depends heavily on the emotions of the person I am interacting with
- A Convolutional Neural Network (CNN) was trained with multiple parameters to identify emotions in images of humans
- The following parametric tuning was then applied to improve accuracy which resulted in overall accuracy of 63%:
 - Increasing the number of Epochs
 - Decreasing batch sizes
 - Increasing the learning rate
- Future considerations for improving this Neural Network include:
 - Testing Gradient Descent with momentum as the optimizer
 - Testing different Loss functions
 - Updating the convolutional layers with varied parameters
- I am passionate about intelligent systems that can improve the life quality of people
- This project serves as a foundation for an ongoing project to build various computer vision and natural language processing functions that can then be used to create a personal assistant device such as a desktop personal assistant or personal robot
- This desktop assistant/personal robot is meant to be a consumer device to help improve the life quality of its customer




Business Problem



- A key problem in the domain of personal assistant/personal robot is to build a deeper connect with it's customer in a way that it enriches the life of the customer
- Current technology solutions are limited to offering helpful services to their customers but no commercial grade device has emotional support features that are persistent
 - There are however some social companion robots, however, they are focussed only on elder care [1]
- Emotion detection serves as a key component of the intelligence that needs to be persistent within the AI based system to create an emotional connection
- Humans innately use facial emotional expressions to infer emotions of others and this guides them to better interact with others
- Extensive research is being actively performed on the various areas of AI application towards emotion recognition [2]
- This project is a first step in the direction of defining algorithm that can focus on facial expressions to identify emotions

[1]: Carley Thornell. 6 AI innovations for those aging in place (Jul 2022). USA Today ([link](#))

[2]: Anvita Saxena, Ashish Khanna, Deepak Gupta (2020). Emotion Recognition and Detection Methods: A Comprehensive Survey. Journal of Artificial Intelligence and Systems, 2, 53–79. ([link](#))



Scope, Approach, Process, and Libraries



Scope:

- Define vision and approach
- Coding of the CNN
- Parametric optimization
- Report and presentation

Approach:

- Understanding existing research on image based emotion recognition and detection
- Determining the appropriate libraries
- Coding framework
- Optimization approach including variable identification and prioritization

Process:

- Setting up the Linux Ubuntu system (GPU drivers), CUDA, and CUDNN
- Installing the key libraries (See below)
- Coding and related troubleshooting
- Identifying baseline performance
- Leveraging prioritized parameters to tune performance

Libraries

- Tensorflow: Tensorflow is an end-to-end platform to develop and deploy machine learning algorithms
- Keras: an open source library for training and deploying machine learning algorithms
- Numpy: library for arrays and mathematical functions
- Pandas: library for reading, cleansing and analyzing data
- Sklearn: library for data processing and visualization



Neural Network & Algorithm Overview



- Convolutional Neural Network (CNN) was chosen for this project as it is the preferred neural networks for small scale image classification [1]
- CNN progressively extracts higher level representations of image contents
- CNN do not require pre-processing and instead uses raw pixels to derive features and infer the object
- Rectified Linear Units (ReLU) was the activation function used to introduce non-linearity into the model
- Adam Optimizer was used for the CNN as it has proven to be better in terms of general performance and parameters requires less tuning
- Using the Convolutional Neural Network (CNN), facial expressions were predicted with an accuracy of 61.7 percent

[1]: Mingxing Tan and Zihang Dai. Toward Fast and Accurate Neural Networks for Image Recognition (Sept '21). Google Research Blog ([link](#))


[2] Akash Ajagekar, Wiki Entry. (Fall '21). Cornell University Computational Optimization Open Textbook ([link](#))



Dataset Overview




- The data used is from the [Kaggle FER2013 collection](#) which has approx 35,9K 48x48 grayscale photos
- The faces in the photos are centered and roughly occupy the same amount of space in each image
- Each of the photos from the dataset represents one of the following emotions :
 - Anger
 - Disgust
 - Fear
 - Happy
 - Neutral
 - Sadness
 - Surprise



CNN & Algorithm Performance Summary (1 of 2)



- The baseline CNN using Adam optimizer performance is 56.9%
- The CNN using SGD optimizer performance is 33.5%
- Due to decreased SGD performance, the Adam Optimizer was selected and then further optimized using parametric tuning
- The Adam Optimizer CNN with highest achieved accuracy is 61,7% which is a 8.4% improvement compared to baseline of 56.9%
- The highest performing CNN has the following key parameters:
 - Batch Size = 16
 - Epoch = 60
 - Learning Rate = 0.002
- **Optimization process key insights**
 - Increasing Epochs from 30 to 75 resulted in accuracy improvements, however, 60 is the max Epoch at which loss rates do not improve
 - Decreasing batch sizes resulted in accuracy improvements; we began with 64 and determined 16 is optimal
 - Adam Learning Rate accuracy is optimal when increased from 0.0001 to 0.0002. However, increasing further results in degradation



CNN & Algorithm Performance Summary (2 of 2)



Base Adam Model Performance Epochs:30,
Batches :64, LR: 0.0001. Accuracy: 56.0%

	precision	recall	f1-score	support
0	0.469	0.463	0.466	501
1	0.000	0.000	0.000	45
2	0.336	0.075	0.123	477
3	0.754	0.872	0.809	921
4	0.405	0.582	0.477	617
5	0.791	0.586	0.673	406
6	0.510	0.601	0.552	622
accuracy			0.569	3589
macro avg	0.466	0.454	0.443	3589
weighted avg	0.551	0.569	0.543	3589

Optimization: Epochs Increased From 30-75
Accuracy: 60.7%

	precision	recall	f1-score	support
0	0.524	0.499	0.511	501
1	0.000	0.000	0.000	45
2	0.462	0.191	0.270	477
3	0.818	0.878	0.847	921
4	0.460	0.478	0.469	617
5	0.762	0.677	0.717	406
6	0.496	0.736	0.593	622
accuracy			0.607	3589
macro avg	0.503	0.494	0.487	3589
weighted avg	0.596	0.607	0.589	3589

Optimization: Batches Decreased From 64
To 16. Accuracy: 61.1%

	precision	recall	f1-score	support
0	0.543	0.491	0.516	501
1	0.667	0.044	0.083	45
2	0.471	0.241	0.319	477
3	0.817	0.848	0.832	921
4	0.501	0.464	0.481	617
5	0.634	0.798	0.707	406
6	0.517	0.707	0.597	622
accuracy			0.611	3589
macro avg	0.593	0.513	0.505	3589
weighted avg	0.604	0.611	0.595	3589


Optimization: LR Increased From
0.0001 to 0.0002. Accuracy: 61.7%

	precision	recall	f1-score	support
0	0.542	0.507	0.524	501
1	0.500	0.022	0.043	45
2	0.472	0.210	0.290	477
3	0.842	0.860	0.851	921
4	0.465	0.499	0.482	617
5	0.719	0.744	0.731	406
6	0.520	0.738	0.610	622
accuracy			0.617	3589
macro avg	0.580	0.511	0.504	3589
weighted avg	0.612	0.617	0.602	3589

Code Overview




- Code has 859 lines in a single Python File and also an associated Jupyter notebook
- Code has five main sections:
 1. **Libraries**
 2. **Data Processing & Setup**
 3. **CNN Model Setup**
 4. **Model Performance, Benchmarking, and Optimization**
- The explanation for each of the four code sections is shared in the next slides



Code Overview: 1. Libraries



- **Tensorflow:** Tensorflow is an end-to-end platform to develop and deploy machine learning algorithms
- **Keras:** an open source library for training and deploying machine learning algorithms
- **Numpy:** library for arrays and mathematical functions
- **Pandas:** library for reading, cleansing and analyzing data
- **Sklearn:** library for data processing and visualization



Code Overview: 2. Data Processing & Setup

- The data used is from the [Kaggle FER2013 collection](#) which has 35887 48x48 grayscale labeled emotion photos
- Below is the chart showing the data distribution by label:



- The data was split into approximately following subsets:

- Training: ~80% (29068, 48, 48, 1)
- Test: ~10% (3589, 48, 48, 1)
- Validation: ~10% (3230, 48, 48, 1)



Code Overview: 3. CNN Model Setup



- The Emotion Recognition and Detection CNN is composed of the following types of layers including their counts:

Layer Type	Overview	Count
Convolutional 2D layer	Performs convolution through scalar multiplication of inputs and weights	6
Batch Normalization	Layer to normalize data between layers using mini batches to speed up training	7
Max Pooling Layer	Takes the max values from batch input and summarizes into a feature map	5
Dropout	Masking layer that nullifies the contribution of some neurons towards the next layer to prevent overfitting	7
Dense	It is the network layer that takes input from all neurons in preceding layer and passing it as input to the next layer	3
Flatten	This layer takes the multiple pooled feature maps and and flattens it into a long vector	1


- Below is the parameter summary for the baseline Convolutional Neural Network (CNN)

```
=====
Total params: 5,810,183
Trainable params: 5,805,191
Non-trainable params: 4,992
=====
```



Code Overview: 4. Model Performance, Benchmarking, and Optimization

- Code runs the CNN Adam model and computes the training and validation accuracy, loss, and confusion matrix
- We then build a Stochastic Gradient Descent Optimization model leveraging the same parameters as the CNN
- Code runs the SGD Model and computes the training and validation accuracy and loss
- We then apply parametric based changes to further optimize the Adam based CNN. Below is the list:
 - Increasing Epochs from 30 to 50
 - Increasing Epochs from 50 to 75
 - Decreasing Batch Size from 64 to 32
 - Decreasing Batch Size from 32 to 16
 - Decreasing Batch Size from 16 to 8
 - Increasing Adam Learning Rate from 0.0001 to 0.0002
 - Increasing Adam Learning Rate from 0.0001 to 0.0009
 - Increasing Adam Learning Rate from 0.0001 to 0.001
 - Increasing Adam Learning Rate from 0.0001 to 0.002



Key Challenges & Learnings For Future




Key Challenges:

- GPU based acceleration ran into several issues:
 - GPU identification by Ubuntu ran into issues due to driver incompatibilities and lack of clear documentation
 - Tensorflow GPU installation using both Nvidia and Tensorflow installations had gaps and resulted in failure to correctly install
 - Even after resolution of Tensorflow-GPU issues using Conda install, Tensorflow caused severe issues resulting in extremely low accuracy of the neural network (~20%)
- The workaround for Tensorflow-GPU issues was to choose the Non-GPU version of Tensorflow, however, due to significant other issues around Ubuntu/Tensorflow, the recommended approach was to use a new system
- Macbook M1 (ARM) was then used to replace the Ubuntu Desktop
- Despite the relatively new nature of Tensorflow on M1 (ARM), it performed well on the Mac but the CNN training time significantly increased

Key Learnings For Future :

- It's critical to not continue investing time into a systems problem that continues to worsen over time (cut losses)
- Always be willing to change an OS/system in sticky situations including using online compute (Colab et al)
- The Ubuntu system is being replaced as it has caused many issues and there is a high probability that these would continue despite clean installs



Future Enhancements & Optimizations

- Implement video camera frame capture to test and validate emotions
- Leverage Gradient Descent with Momentum for Optimization
- Experiment with other Loss Functions especially Dice Coefficient
- Update parameters for the Convolution layers
- Leverage other larger datasets for emotions






Resources



Project Resource Links



- Github repository homepage ([Link](#))
- Jupyter Notebook ([Link](#))
- Python code file ([Link](#))
- Project report ([Link](#))







Appendix




CNN Visualized (1 of 2)




Conv Layer



Max-Pooling Layer



Dense Layer



CNN Visualized (2 of 2)

