

Summary (TL;DR)

- This project leverages Deep Learning to recognize and detect human emotions in images
- Human emotions are fundamental to our experience on our planet, as emotions affect everyone around us [1]
- One of the key inspirations for this project came from interacting with people; I realized that my interaction depends heavily on the emotions of the person I am interacting with
- A Convolutional Neural Network (CNN) was trained with multiple parameters to identify emotions in images of humans
- The following parametric tuning was then applied to improve accuracy which resulted in overall accuracy of 63%:
 - Increasing the number of Epochs
 - Decreasing batch sizes
 - Increasing the learning rate
- Future considerations for improving this Neural Network include:
 - Testing Gradient Descent with momentum as the optimizer
 - Testing different Loss functions
 - Updating the convolutional layers with varied parameters
- I am passionate about intelligent systems that can improve the life quality of people
- This project serves as a foundation for an ongoing project to build various computer vision and natural language processing functions that can then be used to create a personal assistant device such as a desktop personal assistant or personal robot
- This desktop assistant/personal robot is meant to be a consumer device to help improve the life quality of its customer

TABLE OF CONTENTS

Business Problem.....	3
Scope, Approach, Process, and Libraries.....	3
Neural Network & Algorithm Overview.....	4
Dataset Overview.....	4
CNN & Algorithm Performance Summary.....	5
Code Overview & Libraries.....	6
Code Overview: 1. Libraries.....	6
Code Overview: 2. Data Processing & Setup.....	7
Code Overview: 3. CNN Model Setup.....	7
Code Overview: 4. Model Performance, Benchmarking, and Optimization.....	8
Key Challenges & Learnings For Future Exhibits:.....	9

Future Enhancements & Optimizations.....	9
APPENDIX.....	10
CNN Visualized (1 of 3).....	10
CNN Visualized (2 of 3).....	12
CNN Visualized (3 of 3).....	12
EXHIBITS.....	13
Data Shape.....	13
Data Check.....	13
Data Head Validation.....	13
Data Distribution.....	14
Data Population Split.....	14
Data Sample Validation.....	14
Data Model Summary.....	15
Visualizing Model Results.....	15
Base Model (Adam): Confusion Matrix.....	16
Base Model (Adam): Performance Metrics.....	17
Stochastic Gradient Descent (Sgd): Performance Comparison.....	17
Optimization - Parametric Tuning Using Adam & Increasing Epochs From 30 To 50.....	18
Performance Metrics.....	18
Confusion Matrix.....	19
Optimization - Parametric Tuning Using Adam & Increasing Epochs From 30 To 75.....	19
Performance Metrics.....	19
Confusion Matrix.....	20
Optimization - Parametric Tuning Using Adam & Decreasing Batch Size from 64 to 32.....	20
Performance Metrics.....	20
Confusion Matrix.....	21
Optimization - Parametric Tuning Using Adam & Decreasing Batch Size from 32 to 16.....	21
Performance Metrics.....	21
Optimization - Parametric Tuning Using Adam & Decreasing Batch Size from 16 to 8.....	22
Performance Metrics.....	22
Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.0002.....	23
Performance Metrics.....	23
Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.0009.....	24
Performance Metrics.....	24
Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.001.....	25
Performance Metrics.....	25
Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.002.....	26
Performance Metrics.....	26

Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.0025.....	27
Performance Metrics.....	27
Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.0021.....	28
Performance Metrics.....	28

Business Problem

- A key problem in the domain of personal assistant/personal robot is to build a deeper connect with it's customer in a way that it enriches the life of the customer
- Current technology solutions are limited to offering helpful services to their customers but no commercial grade device has emotional support features that are persistent
- There are however some social companion robots, however, they are focussed only on elder care [1]
- Emotion detection serves as a key component of the intelligence that needs to be persistent within the AI based system to create an emotional connection
- Humans innately use facial emotional expressions to infer emotions of others and this guides them to better interact with others
- Extensive research is being actively performed on the various areas of AI application towards emotion recognition [2]
- This project is a first step in the direction of defining algorithm that can focus on facial expressions to identify emotions

[1]: Carley Thornell. 6 AI innovations for those aging in place (Jul 2022). USA Today ([link](#))

[2]: Anvita Saxena, Ashish Khanna, Deepak Gupta (2020). Emotion Recognition and Detection Methods: A Comprehensive Survey. *Journal of Artificial Intelligence and Systems*, 2, 53–79. ([link](#))

Scope, Approach, Process, and Libraries

Scope:

- Define vision and approach
- Coding of the CNN
- Parametric optimization
- Report and presentation

Approach:

- Understanding existing research on image based emotion recognition and

- detection
- Determining the appropriate libraries
- Coding framework
- Optimization approach including variable identification and prioritization

Process:

- Setting up the Linux Ubuntu system (GPU drivers), CUDA, and CUDNN
- Installing the key libraries (See below)
- Coding and related troubleshooting
- Identifying baseline performance
- Leveraging prioritized parameters to tune performance

Libraries

- Tensorflow: Tensorflow is an end-to-end platform to develop and deploy machine learning algorithms
- Keras: an open source library for training and deploying machine learning algorithms
- Numpy: library for arrays and mathematical functions
- Pandas: library for reading, cleansing and analyzing data
- Sklearn: library for data processing and visualization

Neural Network & Algorithm Overview

- Convolutional Neural Network (CNN) was chosen for this project as it is the preferred neural networks for small scale image classification [1]
- CNN progressively extracts higher level representations of image contents
- CNN do not require pre-processing and instead uses raw pixels to derive features and infer the object
- Rectified Linear Units (ReLU) was the activation function used to introduce non-linearity into the model
- Adam Optimizer was used for the CNN as it has proven to be better in terms of general performance and parameters requires less tuning
- Using the Convolutional Neural Network (CNN), facial expressions were predicted with an accuracy of 61.7 percent

[1]: Mingxing Tan and Zihang Dai. Toward Fast and Accurate Neural Networks for Image Recognition (Sept '21). Google Research Blog ([link](#))

[2] Akash Ajagekar, Wiki Entry. (Fall '21). Cornell University Computational Optimization Open Textbook ([link](#))

Dataset Overview

- The data used is from the Kaggle FER2013 collection which has approx 35,9K 48x48 grayscale photos

- The faces in the photos are centered and roughly occupy the same amount of space in each image
- Each of the photos from the dataset represents one of the following emotions :
 - Anger
 - Disgust
 - Fear
 - Happy
 - Neutral
 - Sadness
 - Surprise

CNN & Algorithm Performance Summary

- The baseline CNN using Adam optimizer performance is 56.9%
- The CNN using SGD optimizer performance is 33.5%
- Due to decreased SGD performance, the Adam Optimizer was selected and then further optimized using parametric tuning
- The Adam Optimizer CNN with highest achieved accuracy is 61.7% which is a 8.4% improvement compared to baseline of 56.9%
- The highest performing CNN has the following key parameters:
 - Batch Size = 16
 - Epoch = 60
 - Learning Rate = 0.002
- **Optimization process key insights**
 - Increasing Epochs from 30 to 75 resulted in accuracy improvements, however, 60 is the max Epoch at which loss rates do not improve
 - Decreasing batch sizes resulted in accuracy improvements; we began with 64 and determined 16 is optimal
 - Adam Learning Rate accuracy is optimal when increased from 0.0001 to 0.0002. However, increasing further results in degradation

Base Adam Model Performance Epochs:30, Batches :64, LR: 0.0001. Accuracy: 56.0%

	precision	recall	f1-score	support
0	0.469	0.463	0.466	501
1	0.000	0.000	0.000	45
2	0.336	0.075	0.123	477
3	0.754	0.872	0.809	921
4	0.495	0.582	0.477	617
5	0.791	0.586	0.673	406
6	0.510	0.601	0.552	622
accuracy			0.569	3589
macro avg	0.466	0.454	0.443	3589
weighted avg	0.551	0.569	0.543	3589

Optimization: Epochs Increased From 30-75 Accuracy: 60.7%

	precision	recall	f1-score	support
0	0.524	0.499	0.511	501
1	0.000	0.000	0.000	45
2	0.462	0.191	0.270	477
3	0.818	0.878	0.847	921
4	0.460	0.478	0.469	617
5	0.762	0.677	0.717	406
6	0.496	0.736	0.593	622
accuracy			0.607	3589
macro avg	0.503	0.494	0.487	3589
weighted avg	0.596	0.607	0.589	3589

Optimization: Batches Decreased From 64 To 16. Accuracy: 61.1%

	precision	recall	f1-score	support
0	0.543	0.491	0.516	501
1	0.667	0.044	0.083	45
2	0.471	0.241	0.319	477
3	0.817	0.848	0.832	921
4	0.501	0.464	0.481	617
5	0.634	0.798	0.707	406
6	0.517	0.707	0.597	622
accuracy			0.611	3589
macro avg	0.593	0.513	0.505	3589
weighted avg	0.604	0.611	0.595	3589

Optimization: LR Increased From 0.0001 to 0.0002. Accuracy: 61.7%

	precision	recall	f1-score	support
0	0.542	0.507	0.524	501
1	0.500	0.022	0.043	45
2	0.472	0.210	0.290	477
3	0.842	0.860	0.851	921
4	0.465	0.499	0.482	617
5	0.719	0.744	0.731	406
6	0.520	0.738	0.610	622
accuracy			0.617	3589
macro avg	0.580	0.511	0.504	3589
weighted avg	0.612	0.617	0.602	3589

Code Overview & Libraries

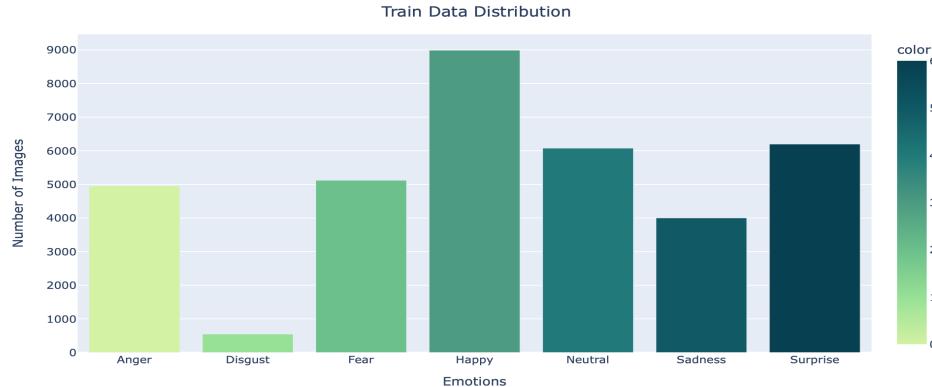
- Code has 859 lines in a single Python File and also an associated Jupyter notebook
- Code has five main sections:
 1. Libraries
 2. Data Processing & Setup
 3. CNN Model Setup
 4. Model Performance, Benchmarking, and Optimization

Code Overview: 1. Libraries

- **Tensorflow:** Tensorflow is an end-to-end platform to develop and deploy machine learning algorithms
- **Keras:** an open source library for training and deploying machine learning algorithms
- **Numpy:** library for arrays and mathematical functions
- **Pandas:** library for reading, cleansing and analyzing data
- **Sklearn:** library for data processing and visualization

Code Overview: 2. Data Processing & Setup

- The data used is from the [Kaggle FER2013 collection](#) which has 35887 48x48 grayscale labeled emotion photos
- Below is the chart showing the data distribution by label:



- The data was split into approximately following subsets:
 - Training: ~80%
 - Test: ~10%
 - Validation: ~10%
 - (29068, 48, 48, 1)
 - (3589, 48, 48, 1)
 - (3230, 48, 48, 1)

Code Overview: 3. CNN Model Setup

- The Emotion Recognition and Detection CNN is composed of the following types of layers including their counts:

Layer Type	Overview	Count
Convolutional 2D layer	Performs convolution through scalar multiplication of inputs and weights	6
Batch Normalization	Layer to normalize data between layers using mini batches to speed up training	7
Max Pooling Layer	Takes the max values from batch input and summarizes into a feature map	5
Dropout	Masking layer that nullifies the contribution of some neurons towards the next layer to prevent overfitting	7

Dense	It is the network layer that takes input from all neurons in preceding layer and passing it as input to the next layer	3
Flatten	This layer takes the multiple pooled feature maps and flattens it into a long vector	1

- Below is the parameter summary for the baseline model Convolutional Neural Network (CNN)
-

Total params: 5,810,183
Trainable params: 5,805,191
Non-trainable params: 4,992

Code Overview: 4. Model Performance, Benchmarking, and Optimization

- Code runs the CNN Adam model and computes the training and validation accuracy, loss, and confusion matrix
- We then build a Stochastic Gradient Descent Optimization model leveraging the same parameters as the CNN
- Code runs the SGD Model and computes the training and validation accuracy and loss
- We then apply parametric based changes to further optimize the Adam based CNN. Below is the list:
 - Increasing Epochs from 30 to 50
 - Increasing Epochs from 50 to 75
 - Decreasing Batch Size from 64 to 32
 - Decreasing Batch Size from 32 to 16
 - Decreasing Batch Size from 16 to 8
 - Increasing Adam Learning Rate from 0.0001 to 0.0002
 - Increasing Adam Learning Rate from 0.0001 to 0.0009

- Increasing Adam Learning Rate from 0.0001 to 0.001
- Increasing Adam Learning Rate from 0.0001 to 0.002

Key Challenges & Learnings For Future Exhibits:

Key Challenges:

- GPU based acceleration ran into several issues:
 - GPU identification by Ubuntu ran into issues due to driver incompatibilities and lack of clear documentation
 - Tensorflow GPU installation using both Nvidia and Tensorflow installations had gaps and resulted in failure to correctly install
 - Even after resolution of Tensorflow-GPU issues using Conda install, Tensorflow caused severe issues resulting in extremely low accuracy of the neural network (~20%)
- The workaround for Tensorflow-GPU issues was to choose the Non-GPU version of Tensorflow, however, due to significant other issues around Ubuntu/Tensorflow, the recommended approach was to use a new system
- Macbook M1 (ARM) was then used to replace the Ubuntu Desktop
- Despite the relatively new nature of Tensorflow on M1 (ARM), it performed well on the Mac but the CNN training time significantly increased

Key Learnings For Future :

- It's critical to not continue investing time into a systems problem that continues to worsen over time (cut losses)
- Always be willing to change an OS/system in sticky situations including using online compute (Colab et al)
- The Ubuntu system is being replaced as it has caused many issues and there is a high probability that these would continue despite clean installs

Future Enhancements & Optimizations

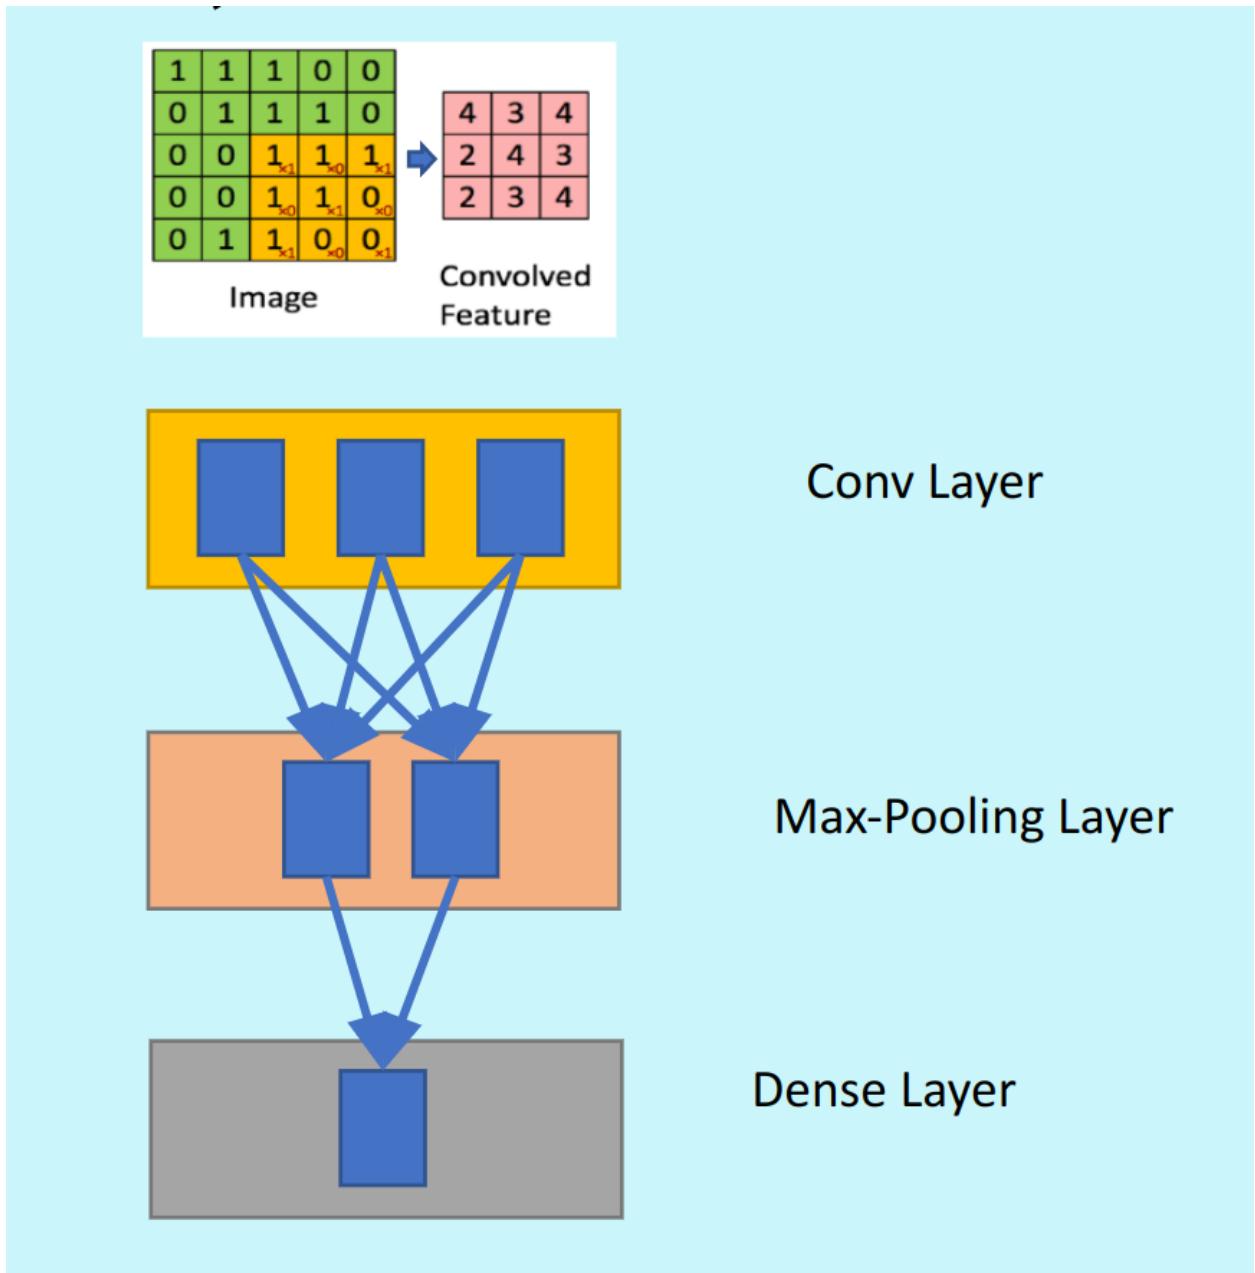
- Implement video camera frame capture to test and validate emotions
- Leverage Gradient Descent with Momentum for Optimization
- Experiment with other Loss Functions especially Dice Coefficient
- Update parameters for the Convolution layers
- Leverage other larger datasets for emotions

Project Resource Links

- Github repository homepage ([Link](#))
- Jupyter Notebook ([Link](#))
- Python code file ([Link](#))
- Project Presentation ([Link](#))

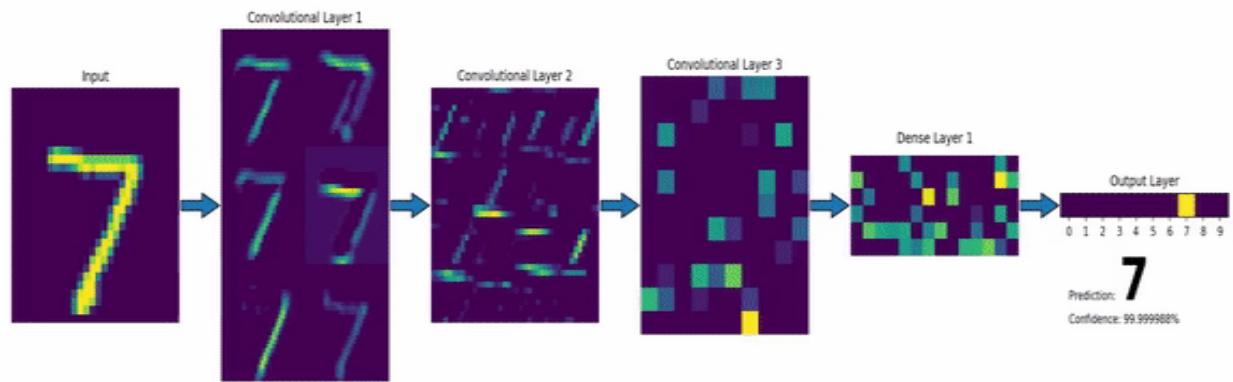
APPENDIX

CNN Visualized (1 of 3)

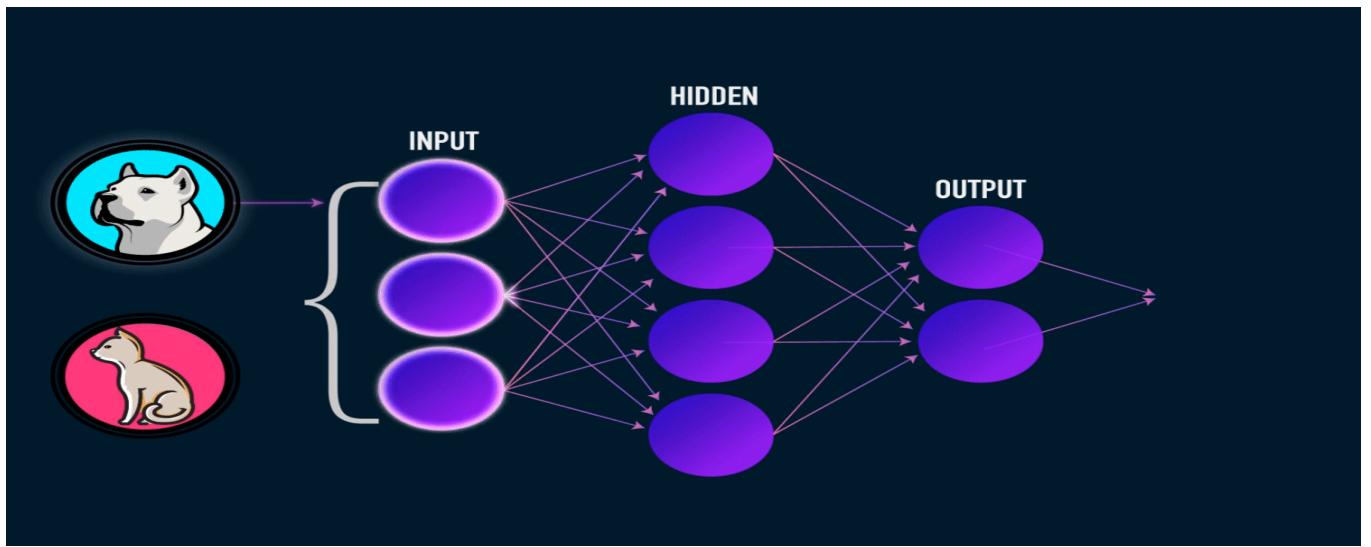


[1]: Mingxing Tan and Zihang Dai. Toward Fast and Accurate Neural Networks for Image Recognition (Sept 2021). Google Research Blog ([link](#))

CNN Visualized (2 of 3)



CNN Visualized (3 of 3)



EXHIBITS

Data Shape

Jupyter Notebook Step 2

(35887, 3)

Data Check

Jupyter Notebook Step 3

```
emotion      0
pixels       0
Usage        0
dtype: int64
```

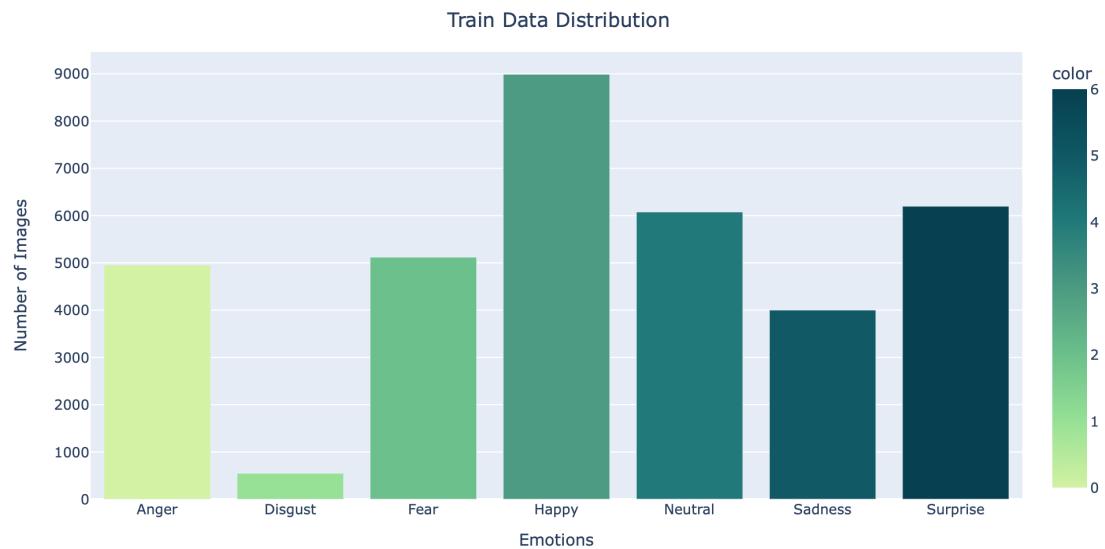
Data Head Validation

Jupyter Notebook Step 4

```
   emotion           pixels      Usage
0      0    70  80  82  72  58  58  60  63  54  58  60  48  89  115  121...
1      0   151  150  147  155  148  133  111  140  170  174  182  15...
2      2   231  212  156  164  174  138  161  173  182  200  106  38...
3      4   24   32   36   30   32   23   19   20   30   41   21   22   32   34   21   1...
4      6   4   0   0   0   0   0   0   0   0   3   15   23   28   48   50   58   84...
```

Data Distribution

Jupyter Notebook Step 5



Data Population Split

Jupyter Notebook Step 10

```
(29068, 48, 48, 1)  
(3589, 48, 48, 1)  
(3230, 48, 48, 1)
```

Data Sample Validation

Jupyter Notebook Step 11



Data Model Summary

Jupyter Notebook Step 13 (Part 1)

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 48, 48, 32)	320
conv2d_1 (Conv2D)	(None, 48, 48, 64)	18496
batch_normalization (BatchN ormalization)	(None, 48, 48, 64)	256
max_pooling2d (MaxPooling2D)	(None, 24, 24, 64)	0
dropout (Dropout)	(None, 24, 24, 64)	0
conv2d_2 (Conv2D)	(None, 24, 24, 128)	204928
batch_normalization_1 (Batch hNormalization)	(None, 24, 24, 128)	512
max_pooling2d_1 (MaxPooling 2D)	(None, 12, 12, 128)	0
dropout_1 (Dropout)	(None, 12, 12, 128)	0
conv2d_3 (Conv2D)	(None, 12, 12, 512)	590336
batch_normalization_2 (Batch hNormalization)	(None, 12, 12, 512)	2048
max_pooling2d_2 (MaxPooling 2D)	(None, 6, 6, 512)	0
dropout_2 (Dropout)	(None, 6, 6, 512)	0
conv2d_4 (Conv2D)	(None, 6, 6, 512)	2359808
batch_normalization_3 (Batch hNormalization)	(None, 6, 6, 512)	2048
max_pooling2d_3 (MaxPooling 2D)	(None, 3, 3, 512)	0
dropout_3 (Dropout)	(None, 3, 3, 512)	0
conv2d_5 (Conv2D)	(None, 3, 3, 512)	2359808
batch_normalization_4 (Batch hNormalization)	(None, 3, 3, 512)	2048
max_pooling2d_4 (MaxPooling 2D)	(None, 1, 1, 512)	0
dropout_4 (Dropout)	(None, 1, 1, 512)	0
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 256)	131328
batch_normalization_5 (Batch hNormalization)	(None, 256)	1024
dropout_5 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 512)	131584
batch_normalization_6 (Batch hNormalization)	(None, 512)	2048
dropout_6 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 7)	3591

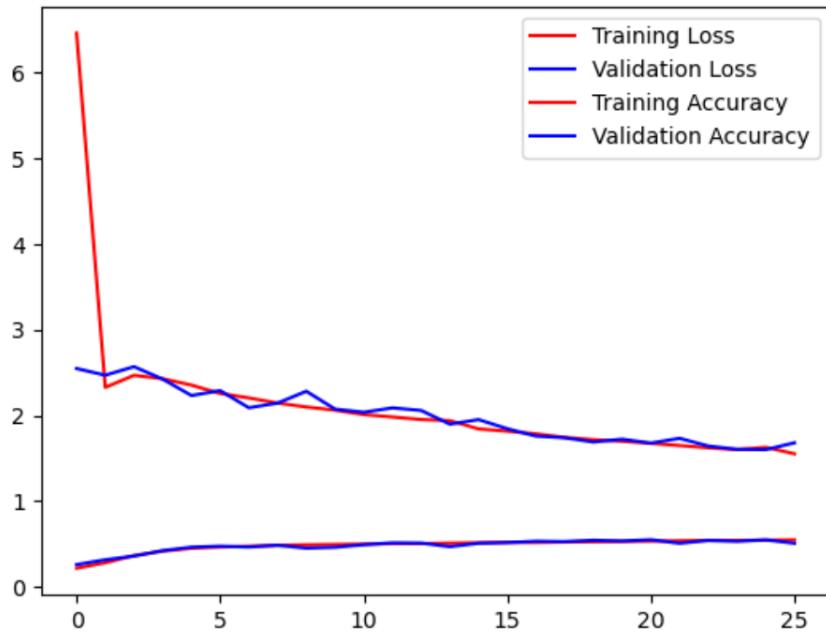
Jupyter Notebook Step 13 (Part 2)

=====

Total params: 5,810,183
Trainable params: 5,805,191
Non-trainable params: 4,992

Visualizing Model Results

Jupyter Notebook Step 15 (Part 1)



Jupyter Notebook Step 15 (Part 2)



Base Model (Adam): Confusion Matrix

Jupyter Notebook Step 16

Confusion Matrix							
	Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise
Actual	232	0	9	50	139	3	68
Anger	28	0	1	4	10	0	2
Disgust	108	0	36	46	173	49	65
Fear	18	0	5	803	36	6	53
Happy	60	0	10	59	359	2	127
Neutral	28	0	38	40	17	238	45
Sadness	21	0	8	63	153	3	374
Surprise	Anger	Disgust	Fear	Happy Predicted	Neutral	Sadness	Surprise

Base Model (Adam): Performance Metrics

Jupyter Notebook Step 17

	precision	recall	f1-score	support
0	0.469	0.463	0.466	501
1	0.000	0.000	0.000	45
2	0.336	0.075	0.123	477
3	0.754	0.872	0.809	921
4	0.405	0.582	0.477	617
5	0.791	0.586	0.673	406
6	0.510	0.601	0.552	622
accuracy			0.569	3589
macro avg	0.466	0.454	0.443	3589
weighted avg	0.551	0.569	0.543	3589

Stochastic Gradient Descent (Sgd): Performance Comparison

Jupyter Notebook Step 18

```
Epoch 30: early stopping
113/113 [=====] - 3s 23ms/step - loss: 9.4962 - accuracy: 0.3349
Test Acc: 0.3349122405052185
```



Optimization - Parametric Tuning Using Adam & Increasing Epochs From 30 To 50

Jupyter Notebook Step 21

Performance Metrics

	precision	recall	f1-score	support
0	0.519	0.547	0.533	501
1	0.000	0.000	0.000	45
2	0.481	0.136	0.212	477
3	0.891	0.754	0.816	921
4	0.390	0.640	0.485	617
5	0.734	0.707	0.720	406
6	0.538	0.643	0.586	622
accuracy			0.589	3589
macro avg	0.508	0.490	0.479	3589
weighted avg	0.608	0.589	0.578	3589

Confusion Matrix

		Confusion Matrix						
		Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise
Actual	Anger	274	0	13	29	115	8	62
	Disgust	28	0	1	1	13	0	2
Happy	Fear	94	0	65	14	203	60	41
	Neutral	16	0	7	694	88	24	92
Neutral	Sadness	80	0	6	12	395	6	118
	Surprise	17	0	36	13	25	287	28
Sadness	Anger	19	0	7	16	174	6	400
	Disgust							
		Happy Predicted				Neutral		
		Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise

Optimization - Parametric Tuning Using Adam & Increasing Epochs From 30 To 75

Jupyter Notebook Step 23

Performance Metrics

	precision	recall	f1-score	support
0	0.524	0.499	0.511	501
1	0.000	0.000	0.000	45
2	0.462	0.191	0.270	477
3	0.818	0.878	0.847	921
4	0.460	0.478	0.469	617
5	0.762	0.677	0.717	406
6	0.496	0.736	0.593	622
accuracy			0.607	3589
macro avg	0.503	0.494	0.487	3589
weighted avg	0.596	0.607	0.589	3589

Confusion Matrix

		Confusion Matrix						
		Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise
Actual	Anger	250	0	21	41	70	10	109
	Disgust	29	0	1	8	4	0	3
Happy	Fear	90	0	91	28	136	52	80
	Neutral	11	0	10	809	31	12	48
Sadness	Neutral	64	0	21	33	295	6	198
	Sadness	17	0	45	24	18	275	27
Surprise	Surprise	16	0	8	46	88	6	458
			Predicted		Happy			

Optimization - Parametric Tuning Using Adam & Decreasing Batch Size from 64 to 32

Jupyter Notebook Step 25

Performance Metrics

	precision	recall	f1-score	support
0	0.509	0.553	0.530	501
1	0.000	0.000	0.000	45
2	0.407	0.300	0.345	477
3	0.838	0.834	0.836	921
4	0.460	0.433	0.446	617
5	0.738	0.702	0.720	406
6	0.544	0.711	0.616	622
accuracy			0.608	3589
macro avg	0.500	0.505	0.499	3589
weighted avg	0.597	0.608	0.599	3589

Confusion Matrix

Confusion Matrix								
	Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise	
Actual	Anger	277	0	49	38	62	9	66
	Disgust	27	0	7	2	7	0	2
	Fear	69	0	143	20	124	63	58
Happy	Anger	24	0	15	768	40	15	59
	Neutral	98	0	55	23	267	8	166
	Sadness	12	0	54	27	9	285	19
	Surprise	37	0	28	38	71	6	442
Predicted								
Anger Disgust Fear Happy Neutral Sadness Surprise								

Optimization - Parametric Tuning Using Adam & Decreasing Batch Size from 32 to 16

Jupyter Notebook Step 27

Performance Metrics

	precision	recall	f1-score	support
0	0.543	0.491	0.516	501
1	0.667	0.044	0.083	45
2	0.471	0.241	0.319	477
3	0.817	0.848	0.832	921
4	0.501	0.464	0.481	617
5	0.634	0.798	0.707	406
6	0.517	0.707	0.597	622
accuracy			0.611	3589
macro avg	0.593	0.513	0.505	3589
weighted avg	0.604	0.611	0.595	3589

Confusion Matrix

Confusion Matrix								
	Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise	
Actual	Anger	246	1	37	43	58	23	93
	Disgust	27	2	6	4	3	2	1
	Fear	69	0	115	27	129	83	54
	Happy	13	0	10	781	25	31	61
	Neutral	66	0	28	37	286	15	185
	Sadness	6	0	36	13	10	324	17
	Surprise	26	0	12	51	60	33	440
Predicted								
Anger Disgust Fear Happy Neutral Sadness Surprise								

Optimization - Parametric Tuning Using Adam & Decreasing Batch Size from 16 to 8

Jupyter Notebook Step 29

Performance Metrics

	precision	recall	f1-score	support
0	0.542	0.507	0.524	501
1	0.500	0.022	0.043	45
2	0.472	0.210	0.290	477
3	0.842	0.860	0.851	921
4	0.465	0.499	0.482	617
5	0.719	0.744	0.731	406
6	0.520	0.738	0.610	622
accuracy			0.617	3589
macro avg	0.580	0.511	0.504	3589
weighted avg	0.612	0.617	0.602	3589

Confusion Matrix

Confusion Matrix								
	Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise	
Actual	Anger	254	1	23	34	79	11	99
	Disgust	27	1	2	2	8	1	4
	Fear	81	0	100	20	149	69	58
	Happy	16	0	7	792	31	15	60
	Neutral	58	0	27	30	308	8	186
	Sadness	17	0	37	20	13	302	17
	Surprise	16	0	16	43	74	14	459
Predicted		Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise

Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.0002

Jupyter Notebook Step 30

Performance Metrics

	precision	recall	f1-score	support
0	0.542	0.507	0.524	501
1	0.500	0.022	0.043	45
2	0.472	0.210	0.290	477
3	0.842	0.860	0.851	921
4	0.465	0.499	0.482	617
5	0.719	0.744	0.731	406
6	0.520	0.738	0.610	622
accuracy			0.617	3589
macro avg	0.580	0.511	0.504	3589
weighted avg	0.612	0.617	0.602	3589

Confusion Matrix

		Confusion Matrix						
		Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise
Actual	Anger	254	1	23	34	79	11	99
	Disgust	27	1	2	2	8	1	4
Happy	Fear	81	0	100	20	149	69	58
	Neutral	16	0	7	792	31	15	60
Sadness	Surprise	58	0	27	30	308	8	186
	Anger	17	0	37	20	13	302	17
Surprise	Disgust	16	0	16	43	74	14	459
	Fear							
		Predicted				Happy		

Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.0009

Jupyter Notebook Step 33

Performance Metrics

	precision	recall	f1-score	support
0	0.543	0.491	0.516	501
1	0.667	0.044	0.083	45
2	0.471	0.241	0.319	477
3	0.817	0.848	0.832	921
4	0.501	0.464	0.481	617
5	0.634	0.798	0.707	406
6	0.517	0.707	0.597	622
accuracy			0.611	3589
macro avg	0.593	0.513	0.505	3589
weighted avg	0.604	0.611	0.595	3589

Confusion Matrix

Confusion Matrix								
	Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise	
Actual	Anger	246	1	37	43	58	23	93
	Disgust	27	2	6	4	3	2	1
	Fear	69	0	115	27	129	83	54
	Happy	13	0	10	781	25	31	61
	Neutral	66	0	28	37	286	15	185
	Sadness	6	0	36	13	10	324	17
	Surprise	26	0	12	51	60	33	440
Predicted								
Anger Disgust Fear Happy Neutral Sadness Surprise								

Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.001

Jupyter Notebook Step 36

Performance Metrics

	precision	recall	f1-score	support
0	0.509	0.487	0.498	501
1	0.000	0.000	0.000	45
2	0.381	0.224	0.282	477
3	0.811	0.857	0.833	921
4	0.468	0.522	0.493	617
5	0.703	0.746	0.724	406
6	0.564	0.669	0.612	622
accuracy			0.608	3589
macro avg	0.491	0.501	0.492	3589
weighted avg	0.588	0.608	0.594	3589

Confusion Matrix

Confusion Matrix							
	Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise
Actual	244	0	40	43	88	17	69
Anger	27	0	5	3	9	0	1
Disgust	79	0	107	29	141	73	48
Fear	19	0	15	789	26	16	56
Happy	73	0	39	42	322	5	136
Neutral	9	0	53	20	10	303	11
Sadness	28	0	22	47	92	17	416
Surprise	Anger	Disgust	Fear	Happy Predicted	Neutral	Sadness	Surprise

Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.002

Jupyter Notebook Step 38

Performance Metrics

	precision	recall	f1-score	support
0	0.509	0.487	0.498	501
1	0.000	0.000	0.000	45
2	0.381	0.224	0.282	477
3	0.811	0.857	0.833	921
4	0.468	0.522	0.493	617
5	0.703	0.746	0.724	406
6	0.564	0.669	0.612	622
accuracy			0.608	3589
macro avg	0.491	0.501	0.492	3589
weighted avg	0.588	0.608	0.594	3589

Confusion Matrix

Confusion Matrix								
	Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise	
Actual	Anger	244	0	40	43	88	17	69
	Disgust	27	0	5	3	9	0	1
	Fear	79	0	107	29	141	73	48
Happy	Anger	19	0	15	789	26	16	56
	Neutral	73	0	39	42	322	5	136
	Sadness	9	0	53	20	10	303	11
	Surprise	28	0	22	47	92	17	416
Predicted								
Anger Disgust Fear Happy Neutral Sadness Surprise								

Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.0025

Jupyter Notebook Step 40

Performance Metrics

	precision	recall	f1-score	support
0	0.509	0.487	0.498	501
1	0.000	0.000	0.000	45
2	0.381	0.224	0.282	477
3	0.811	0.857	0.833	921
4	0.468	0.522	0.493	617
5	0.703	0.746	0.724	406
6	0.564	0.669	0.612	622
accuracy			0.608	3589
macro avg	0.491	0.501	0.492	3589
weighted avg	0.588	0.608	0.594	3589

Confusion Matrix

Confusion Matrix								
	Anger	Disgust	Fear	Happy	Neutral	Sadness	Surprise	
Actual	Anger	244	0	40	43	88	17	69
	Disgust	27	0	5	3	9	0	1
	Fear	79	0	107	29	141	73	48
Happy	Anger	19	0	15	789	26	16	56
	Neutral	73	0	39	42	322	5	136
	Sadness	9	0	53	20	10	303	11
	Surprise	28	0	22	47	92	17	416
Predicted								

Optimization - Parametric Tuning Using Adam & Increasing Learning Rate (LR) from 0.0001 to 0.0021

Jupyter Notebook Step

Performance Metrics

	precision	recall	f1-score	support
0	0.509	0.487	0.498	501
1	0.000	0.000	0.000	45
2	0.381	0.224	0.282	477
3	0.811	0.857	0.833	921
4	0.468	0.522	0.493	617
5	0.703	0.746	0.724	406
6	0.564	0.669	0.612	622
accuracy			0.608	3589
macro avg	0.491	0.501	0.492	3589
weighted avg	0.588	0.608	0.594	3589

Confusion Matrix

		Confusion Matrix													
		Anger		Disgust		Fear		Happy		Neutral		Sadness		Surprise	
Actual	Happy	Anger		Disgust		Fear		Happy		Neutral		Sadness		Surprise	
		244	0	40	43	88	17	69	789	26	16	56	416	136	
Anger		244	0	40	43	88	17	69							
Disgust		27	0	5	3	9	0	1							
Fear		79	0	107	29	141	73	48							
Happy		19	0	15	789	26	16	56							
Neutral		73	0	39	42	322	5	136							
Sadness		9	0	53	20	10	303	11							
Surprise		28	0	22	47	92	17								
	Anger		Disgust		Fear		Happy		Predicted		Neutral		Sadness		Surprise